

# Lab1\_Block2\_Assignment2

2022-12-06

## Assignment 2. Mixture models

The EM algorithm was implemented for the Bernoulli mixture model. Note that the R code was adapted slightly from the provided template to produce output that better aligned with the R markdown format (though it produces the same results as using the provided template from the beginning for each value of  $M$ ).

The Bernoulli mixture model is:

$$p(\mathbf{x}) = \sum_{m=1}^M \pi_m \text{Bern}(\mathbf{x}|\boldsymbol{\mu}_m)$$

where  $\mathbf{x} = \{x_1, x_2, \dots, x_D\}$  is a  $D$ -dimensional binary random vector,  $\pi_m = p(y = m)$  and

$$\text{Bern}(\mathbf{x}|\boldsymbol{\mu}_m) = p(\mathbf{x}|y = m) = \prod_{d=1}^D \mu_{m,d}^{x_d} (1 - \mu_{m,d})^{(1-x_d)}$$

where  $\boldsymbol{\mu} = (\mu_{m,1}, \dots, \mu_{m,D})$  is a  $D$ -dimensional vector of probabilities.

The log likelihood of the dataset  $\{\mathbf{x}_i^n\}$  is:

$$\sum_{i=1}^n \log p(\mathbf{x}_i)$$

The weights in the EM algorithm are calculated as  $p(y = m|\mathbf{x}_i, \hat{\boldsymbol{\mu}})$

We have that  $p(\mathbf{x}, y) = p(\mathbf{x}|y)p(y)$  and that  $p(\mathbf{x}, y) = p(y|\mathbf{x})p(\mathbf{x})$ , thus

$$p(y|\mathbf{x}) = \frac{p(y)p(\mathbf{x}|y)}{p(\mathbf{x})}$$

The weights then become:


$$p(y = m|\mathbf{x}_i) = \frac{p(y = m)p(\mathbf{x}_i|y = m)}{p(\mathbf{x}_i)} = \frac{\hat{\pi}_m \text{Bern}(\mathbf{x}_i|\hat{\boldsymbol{\mu}}_m)}{\sum_{m=1}^M \hat{\pi}_m \text{Bern}(\mathbf{x}_i|\hat{\boldsymbol{\mu}}_m)}$$

Finally, the parameters are updated according to:

$$\hat{\pi}_m = \frac{1}{n} \sum_{i=1}^n w_i(m),$$
$$\hat{\boldsymbol{\mu}}_m = \frac{1}{\sum_{i=1}^n w_i(m)} \sum_{i=1}^n w_i(m) \mathbf{x}_i$$

The true parameter values are displayed below.

True  $\mu$ :




True  $\pi$ :

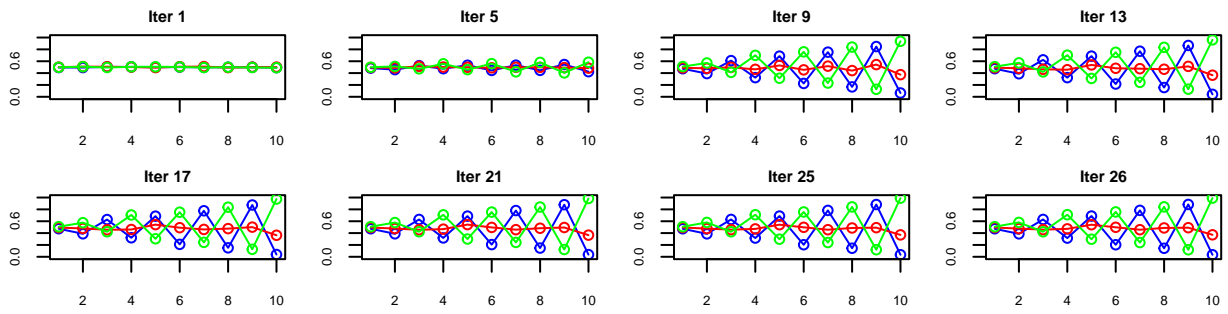
```
#> [1] 0.3333333 0.3333333 0.3333333
```

The EM algorithm was run for  $M = 2, 3$  and 4. Below, estimates of  $\mu$  are displayed for the first, then every fourth iteration as well as after the final iteration. For overview and completeness, the final  $\mu$  are also displayed after the iteration plots, along with the estimated values for  $\pi$ .

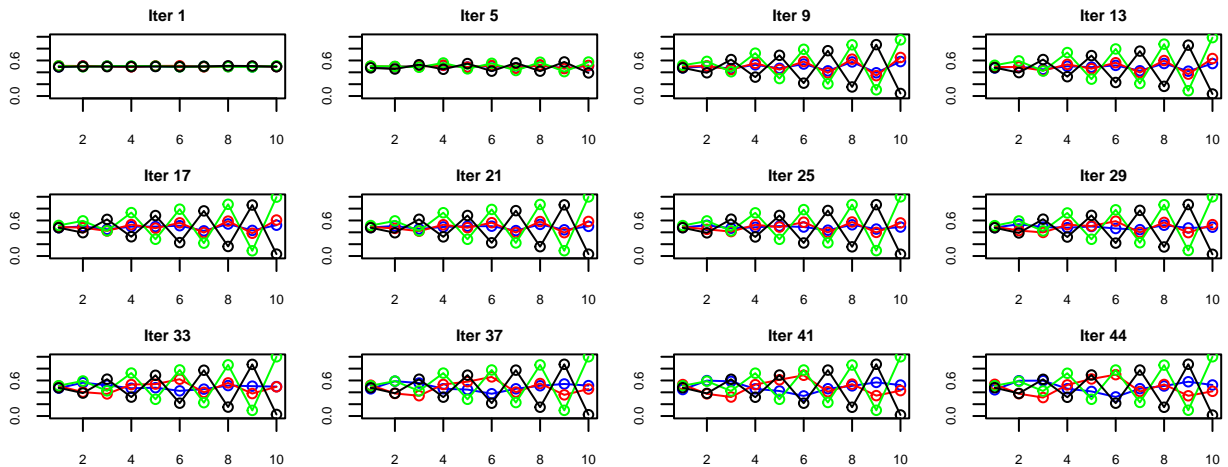
#> M = 2, start values of pi (rounded):0.49921, 0.50079



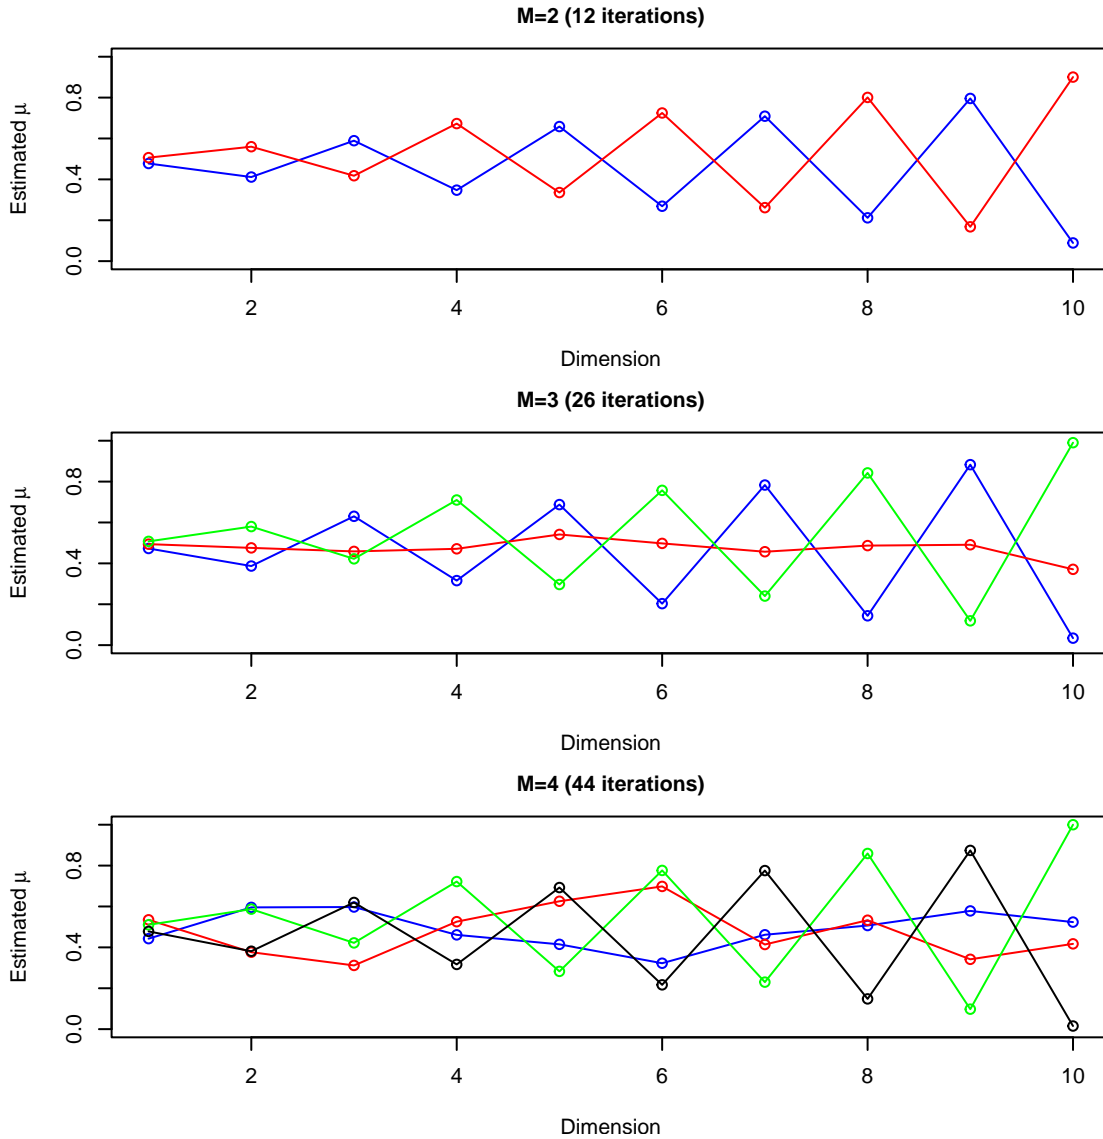
#> M = 3, start values of pi (rounded):0.33261, 0.33366, 0.33374



#> M = 4, start values of pi (rounded):0.24918, 0.24997, 0.25003, 0.25082



The values of the estimated  $\mu$  and  $\pi$  are displayed below:



Estimated  $\pi$ :

```
#> Estimated pi M=2:
```

```
#> 0.497125 0.502875
```


```
#> Estimated pi M=3:
```

```
#> 0.3416794 0.2690298 0.3892909
```

```
#> Estimated pi M=4:
```

```
#> 0.1547196 0.1418652 0.3514089 0.3520062
```

The log likelihood was also plotted versus the iteration number:



It can be seen that the number of iterations of the EM algorithm increased with increasing  $M$ , and that the maximum log likelihood was obtained for  $M = 4$ , even though the true  $M = 3$ . This can be explained by that the higher the number of clusters ( $M$ ), the better fit can be obtained to the data.

## Appendix: All code for this report

```
knitr::opts_chunk$set(comment = "#>",
                        echo = FALSE)

#-----#
#-----#
#- Assignment 2
#-----#
#-----#

#-----#
# Functions
#-----#

#- Function to calculate  $p(x/y)$  for given class  $m$ , with param  $\mu$ 
#  $\mu$  is  $D$ -dimensional vector
#  $x$  is matrix of dimension  $n \times D$ 
bern <- function(x, mu) {
  D <- dim(x)[2]
  stopifnot(length(mu)==D)
  n <- dim(x)[1]
  res <- numeric(n)
  for (i in 1:n) {
    prod_d <- 1
    for (d in 1:D) {
      term_d <- (mu[d]^x[i,d])*((1-mu[d])^(1-x[i,d]))
      prod_d <- term_d*prod_d
    }
    res[i] <- prod_d
  }
  return(res)
}

#- Function to calculate  $p(x)$ 
#  $\mu_v$  - vector of  $\mu$ 
#  $\pi_v$  -vector with  $\pi$ 
px <- function(x, mu_v, pi_v) {
  M <- dim(mu_v)[1]
  sum_px <- 0
  for (m in 1:M) {
    sum_px <- sum_px + (bern(x, mu_v[m,])*pi_v[m])
  }
  return(sum_px)
}

#-----#
# Main program
#-----#

#set.seed(1234567890)

max_it <- 100 # max number of EM iterations
min_change <- 0.1 # min change in log lik between two consecutive iterations
n=1000 # number of training points
```

```

D=10 # number of dimensions
x <- matrix(nrow=n, ncol=D) # training data

true_pi <- vector(length = 3) # true mixing coefficients
true_mu <- matrix(nrow=3, ncol=D) # true conditional distributions
true_pi=c(1/3, 1/3, 1/3)
true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
plot(true_mu[1,], type="o", col="blue", ylim=c(0,1), xlab = "Dimension",
      ylab=expression(mu), main=expression(paste("True ",mu)))
points(true_mu[2,], type="o", col="red")
points(true_mu[3,], type="o", col="green")

#- Note: this is done in loop to obtain same seed as would have been obtained
# when running the original script separately for each M
# #- Producing the training data
# for(i in 1:n) {
#   m <- sample(1:3,1,prob=true_pi)
#   for(d in 1:D) {
#     x[i,d] <- rbinom(1,1,true_mu[m,d])
#   }
# }

true_pi

#M=3 # number of clusters
llik <- matrix(nrow=max_it, ncol=3) # log likelihood of the EM iterations
mu_list <- list() # save final mu for each M
pi_list <- list() # save final pi for each M
iter_list <- list() # save number of iter for each M

for (M in 2:4) {

  cat(paste0("M = ",M))
  set.seed(1234567890)
  #- Producing the training data
  for(i in 1:n) {
    m <- sample(1:3,1,prob=true_pi)
    for(d in 1:D) {
      x[i,d] <- rbinom(1,1,true_mu[m,d])
    }
  }

  w <- matrix(nrow=n, ncol=M) # weights
  pi <- vector(length = M) # mixing coefficients
  mu <- matrix(nrow=M, ncol=D) # conditional distributions
  #llik <- vector(length = max_it) # log likelihood of the EM iterations

  #- Random initialization of the parameters
  pi <- runif(M,0.49,0.51)
  pi <- pi / sum(pi)
  for(m in 1:M) {

```

```

    mu[m,] <- runif(D,0.49,0.51)
}
#pi
#mu
pi_str <- toString(round(pi,5))
cat(paste0(", start values of pi (rounded):", pi_str))
par(mfrow=c(3,4),mai = c(0.3, 0.25, 0.2, 0.15))
for(it in 1:max_it) {
  #- Plot 1st then each 4th iter

  if(it==1 | ((it-1)%4==0)) {
    plot(mu[1,], type="o", col="blue", ylim=c(0,1), main=paste0("Iter ", it),
         cex.main=0.8, cex.axis=0.6)
    points(mu[2,], type="o", col="red")
    if (dim(mu)[1] >= 3) {
      points(mu[3,], type="o", col="green")
    }
    if (dim(mu)[1] == 4) {
      points(mu[4,], type="o", col="black")
    }
  }
  Sys.sleep(0.5)

  #- E-step: Computation of the weights
  #####
  #- Your code here
  sum_px <- px(x=x, mu_v=mu, pi_v=pi)
  for (m in 1:M) {
    w[,m] <- (pi[m]*bern(x, mu[m,]))/sum_px
  }
  #####

  #- Log likelihood computation.
  #####
  #- Your code here
  llik[it,M-1] <- sum(log(sum_px))

  #####

  #cat("iteration: ", it, "log likelihood: ", llik[it,M-1], "\n")
  #flush.console()

  #- Stop if the log likelihood has not changed significantly
  #####
  #- Your code here
  # The value of the observed data log-likelihood increases at each iteration of
  # the procedure, unless it has reached a stationary point
  if (it > 1) {
    llik_ch <- llik[it,M-1]-llik[it-1,M-1]
    if (llik_ch < min_change) {
      break
    }
  }
}

```



```
#####

#- M-step: ML parameter estimation from the data and weights
#####
#- Your code here
for (m in 1:M) {
  pi[m] <- (1/n)*sum(w[,m])
  mu[m,] <- (1/sum(w[,m]))*colSums(w[,m]*x)
}
#####

# - Plot likelihood vs iterations
#plot(llik[1:it,M-1], type="o")

#- Plot last iteration

plot(mu[1,], type="o", col="blue", ylim=c(0,1), main=paste0("Iter ", it),
      cex.main=0.8, cex.axis=0.6)
points(mu[2,], type="o", col="red")
if (dim(mu)[1] >= 3) {
  points(mu[3,], type="o", col="green")
}
if (dim(mu)[1] == 4) {
  points(mu[4,], type="o", col="black")
}

#pi
#mu

mu_list[[M-1]] <- mu
pi_list[[M-1]] <- pi
iter_list[[M-1]] <- it
}

# - Plot final parameters
par(mfrow=c(3,1), mai=c(0.5,0.8,0.35,0.5))
for (M in 2:4) {
  mu <- mu_list[[M-1]]
  it <- iter_list[[M-1]]
  plot(mu[1,], type="o", col="blue", ylim=c(0,1), xlab="Dimension",
        ylab=expression(paste("Estimated ", mu)),
        main=paste0("M=", M, " (", it, " iterations)"), cex.main=1)
  points(mu[2,], type="o", col="red")
  if (dim(mu)[1] >= 3) {
    points(mu[3,], type="o", col="green")
  }
  if (dim(mu)[1] == 4) {
    points(mu[4,], type="o", col="black")
  }
}
}
```

```

cat("Estimated pi M=2:\n")
cat(pi_list[[1]], "\n")

cat("Estimated pi M=3:\n")
cat(pi_list[[2]], "\n")

cat("Estimated pi M=4:\n")
cat(pi_list[[3]], "\n")

# - Plot likelihood vs iterations
par(mfrow=c(3,1), mai=c(0.3,1,0.3,1))
plot(llik[1:it,1], type="o", xlab="Iteration", ylab="Log likelihood",
     main=paste("M=2, final log likelihood=",
                round(llik[iter_list[[1]],1],1), ", ",
                iter_list[[1]], "iterations"))
plot(llik[1:it,2], type="o", xlab="Iteration", ylab="Log likelihood",
     main=paste("M=3, final log likelihood=",
                round(llik[iter_list[[2]],2],1), ", ",
                iter_list[[2]], "iterations"))
plot(llik[1:it,3], type="o", xlab="Iteration", ylab="Log likelihood",
     main=paste("M=4, final log likelihood=",
                round(llik[iter_list[[3]],3],1), ", ",
                iter_list[[3]], "iterations"))

```